Unsupervised and Unstructured Machine Learning

BA820 – Mohannad Elhamod

Intro to Clustering

Cluster Analysis

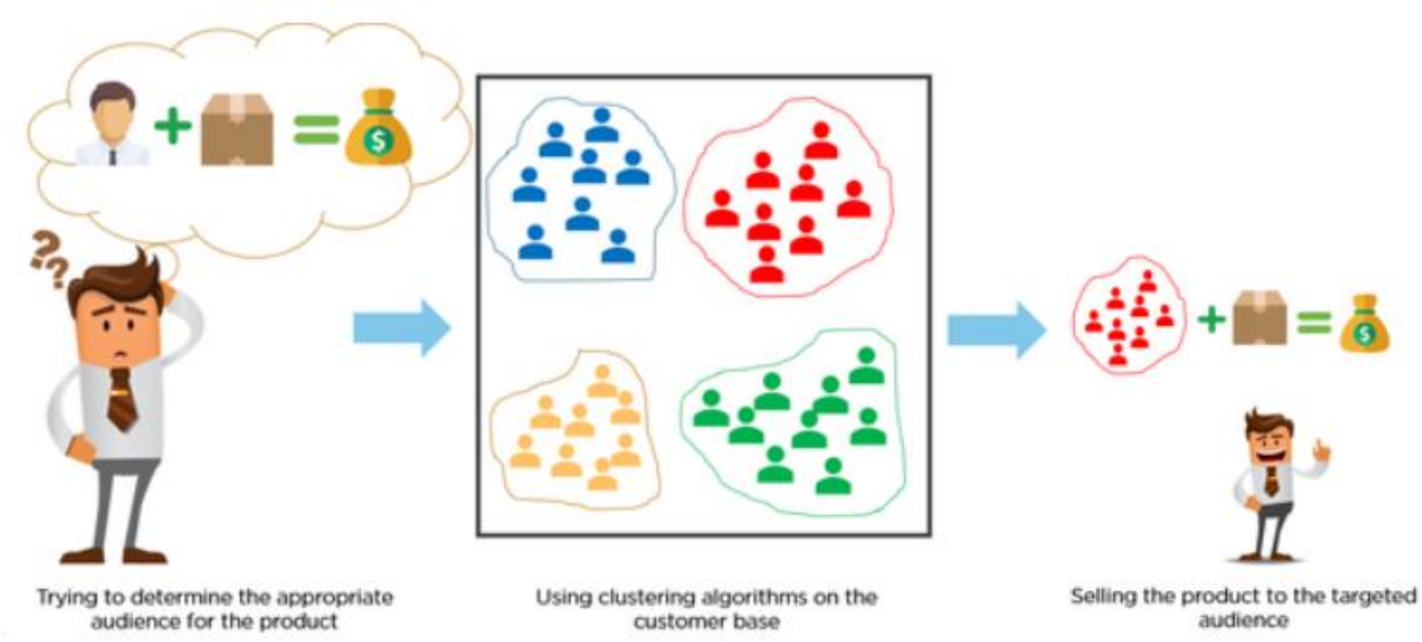

medium.com

What is Cluster Analysis?

- Placing objects in groups such that:
  - the objects in a group are similar (or related) to one another.
  - They are different from (or unrelated to) the objects in other groups.
- We need a (metric/measure/objective function) to measure the (distance/similarity) of the (objects/clusters).

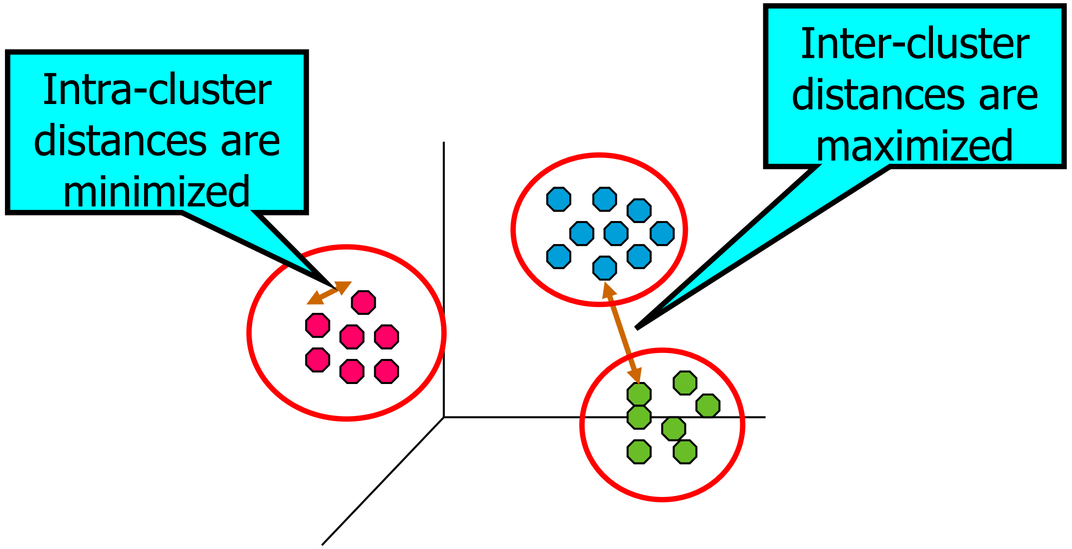

Clusters are in the eye of the beholder

How many clusters?

Six Clusters

Two Clusters

Four Clusters

Clusters come in all shapes and sizes

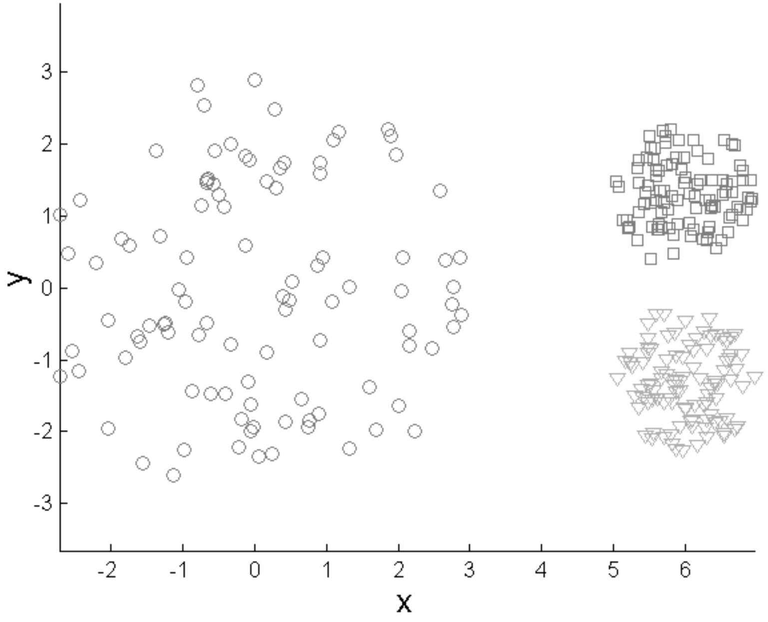

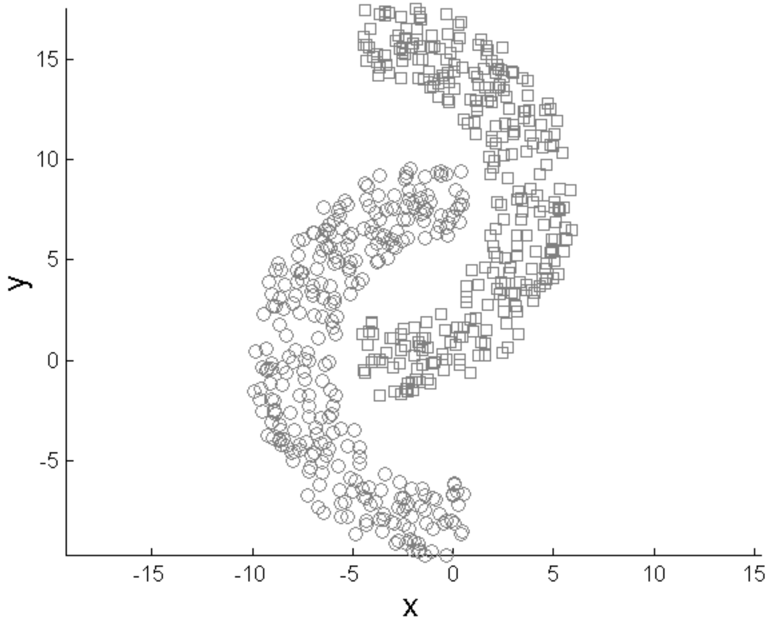

Hierarchical Clustering

BOSTON UNIVERSITY

Boston University Questrom School of Business

Hierarchical Clustering

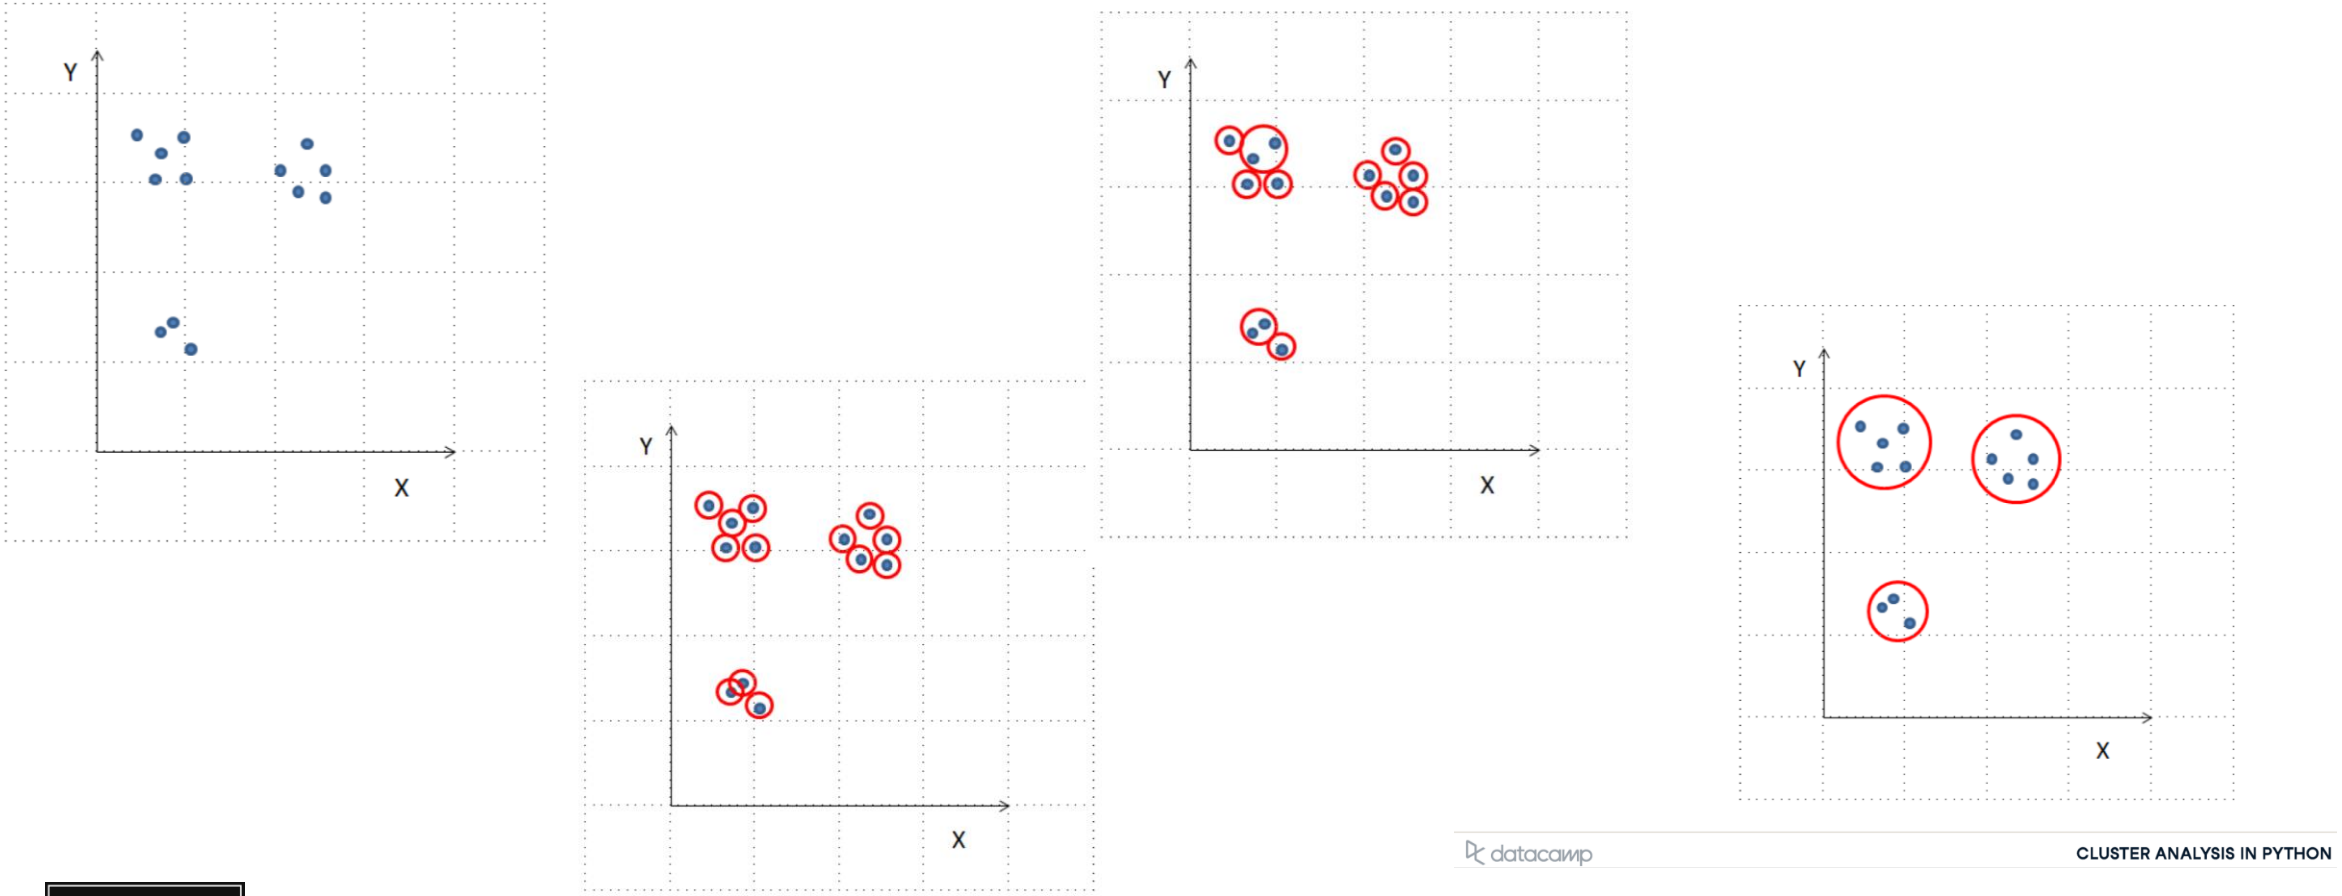

Hierarchical Clustering

- Produces a set of nested clusters organized as a hierarchical tree
- Can be visualized as a dendrogram
  - A tree like diagram that records the sequences of merges or splits.

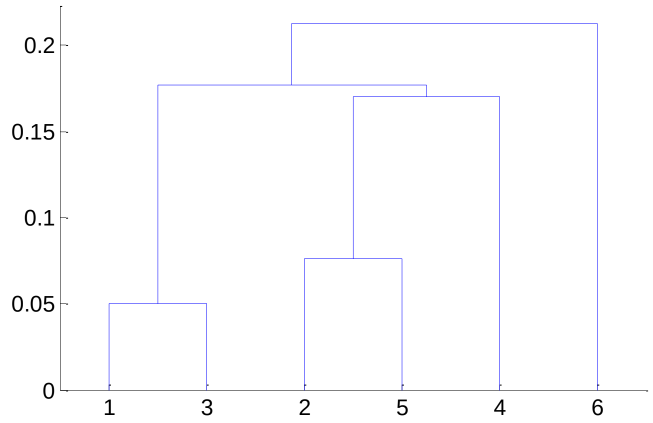

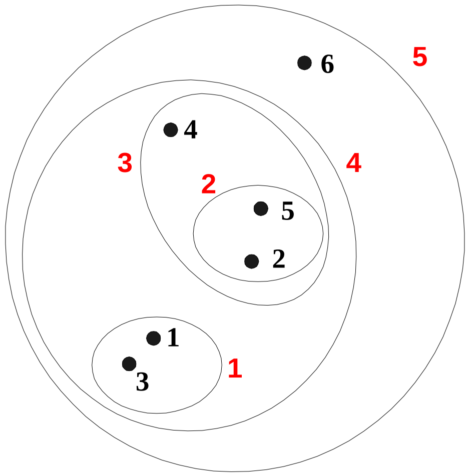

Hierarchical Clustering

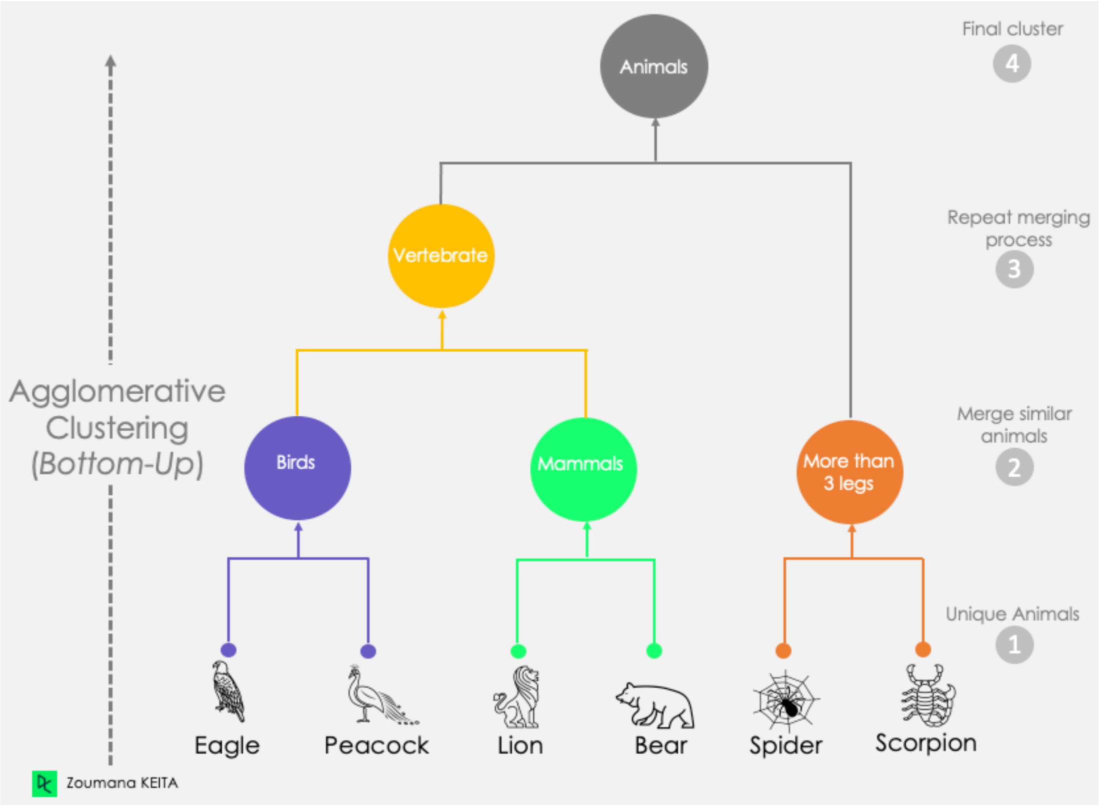

Hierarchical Clustering

- **Key Idea: Successively merge closest clusters**
- Basic algorithm
  1. Compute the proximity matrix
  2. Let each data point be a cluster
  3. **Repeat**
  4. Merge the two closest clusters
  5. Update the proximity matrix
  6. **Until** only a single cluster remains
- Key operation is the computation of the proximity of two clusters
  - Different approaches to defining the distance between clusters distinguish the different algorithms

Steps 1 and 2

- Start with clusters of individual points and a proximity matrix

	p1	p2	p3	p4	p5	. . .
p1						
p2						
p3						
p4						
p5						
.						
.						
.						

Proximity Matrix

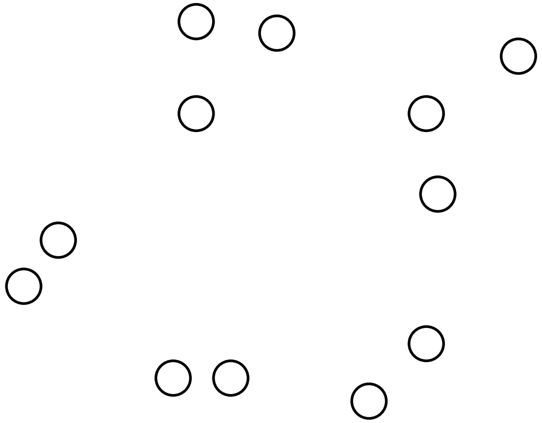

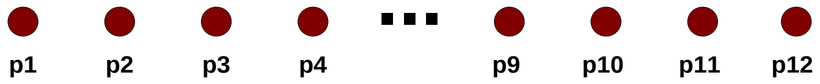

Intermediate Situation

After some merging steps, we have some clusters

	C1	C2	C3	C4	C5
C1					
C2					
C3					
C4					
C5					

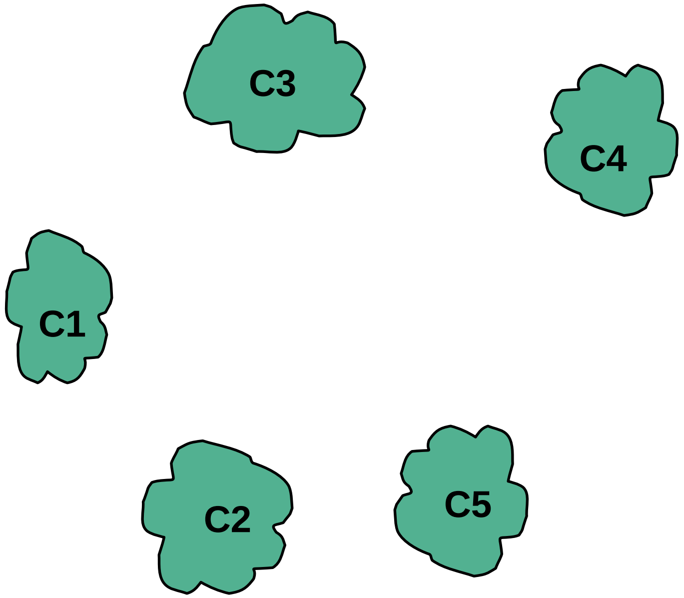

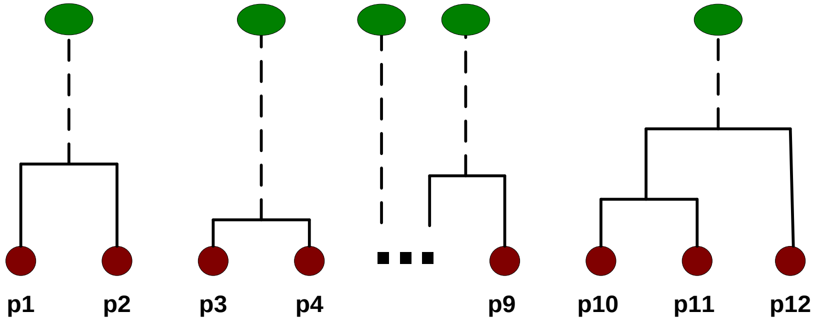

Step 4

We want to merge the two closest clusters (C2 and C5) and update the proximity matrix.

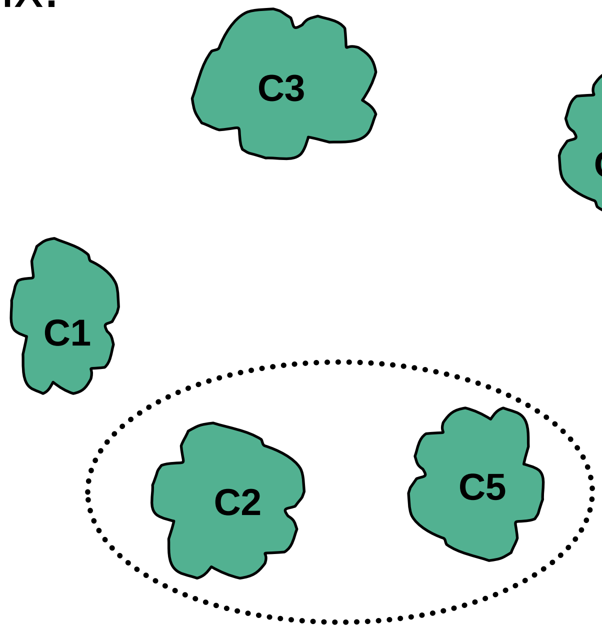

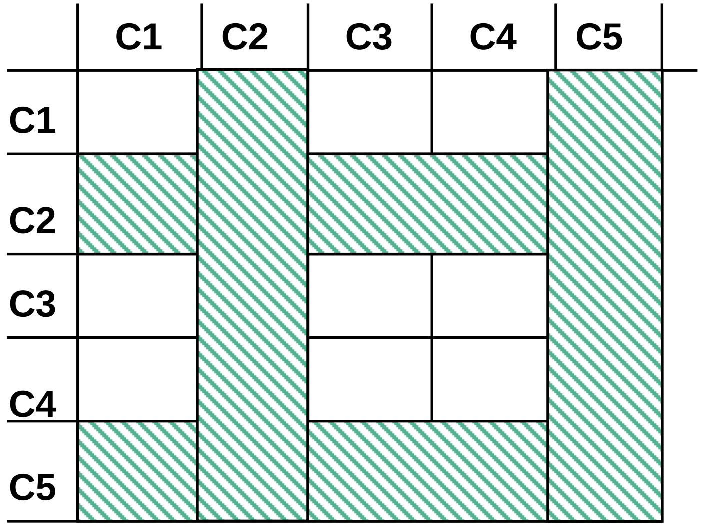

Proximity Matrix

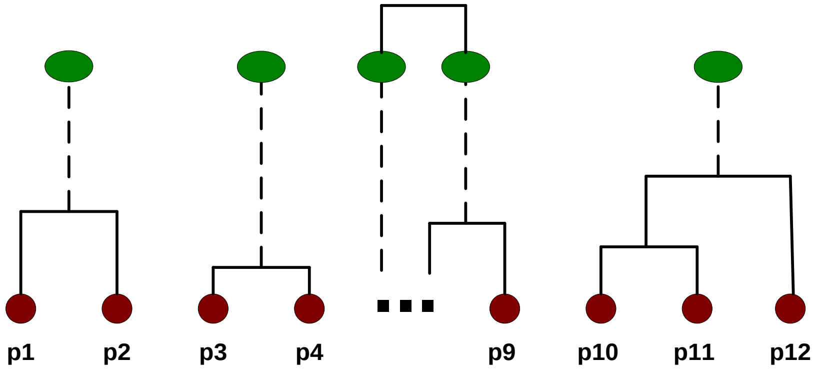

Step 5

The question is "How do we update the proximity matrix?"

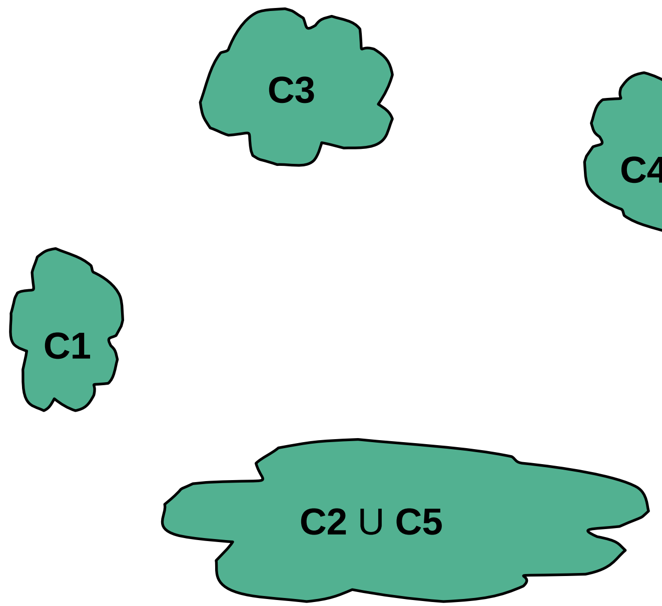

	C1	C2 U C5	C3	C4
C1		?		
C2 U C5	?	?	?	?
C3		?		
C4		?		

Proximity Matrix

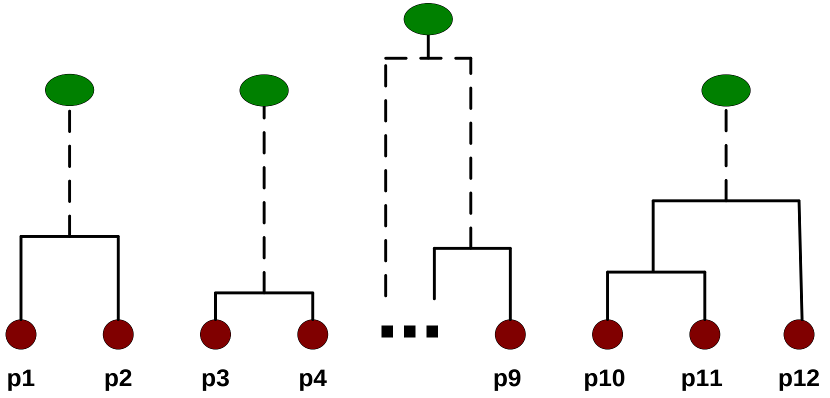

How to Define Inter-Cluster Distance

- MIN (Single Link)
- MAX (Complete Linkage)
- Group Average
  - Ward's Method uses squared error
- Distance Between Centroids

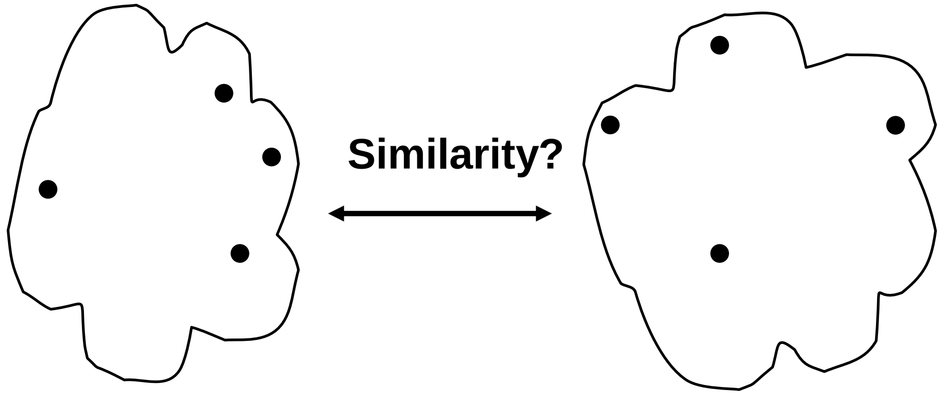

How to Define Inter-Cluster Similarity

- MIN (Single Link)
  - Sensitive to noise.

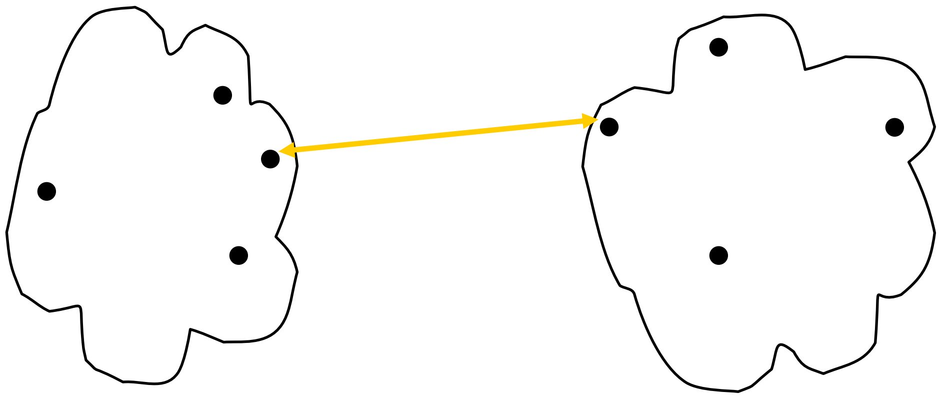

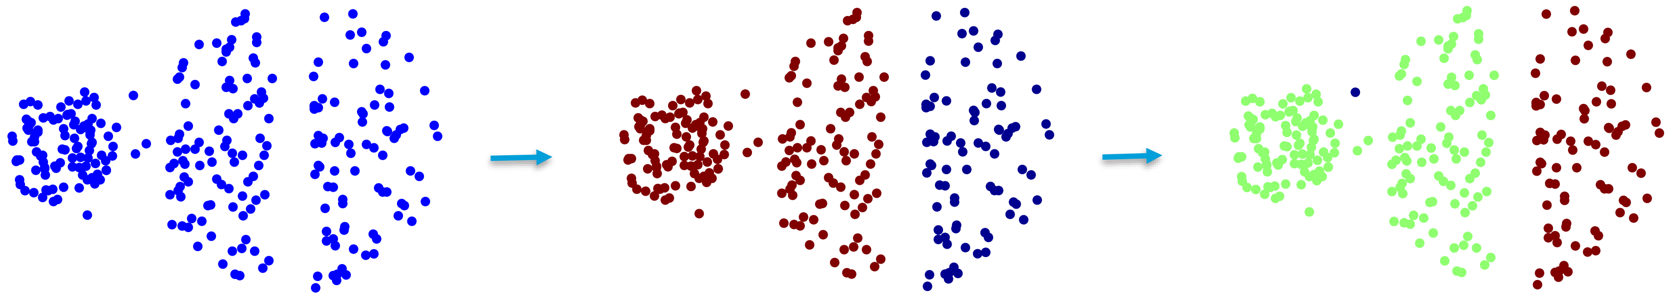

How to Define Inter-Cluster Similarity

- MAX (Complete Linkage)
  - Less susceptible to noise.
  - Breaks larger clusters.

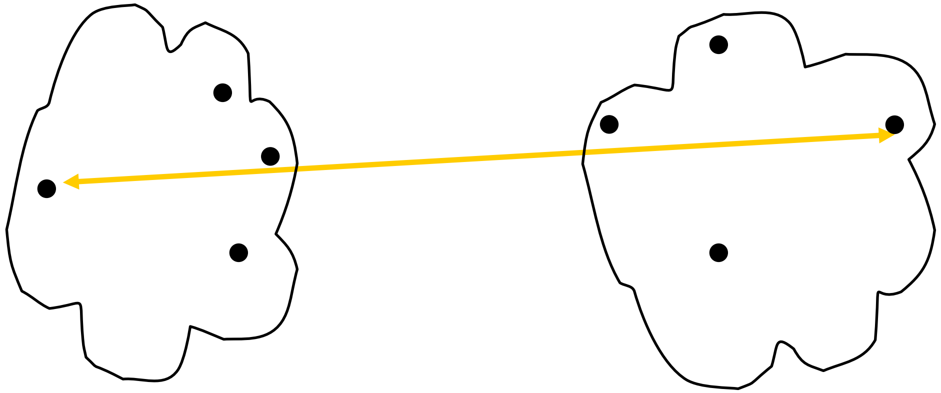

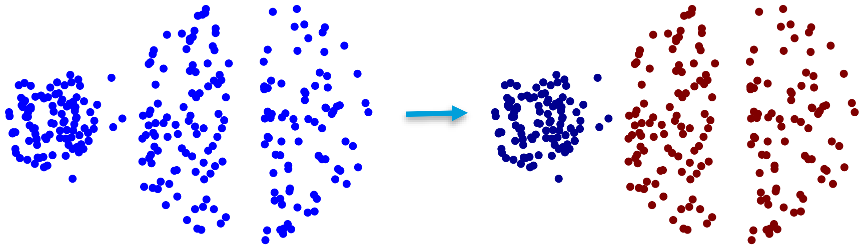

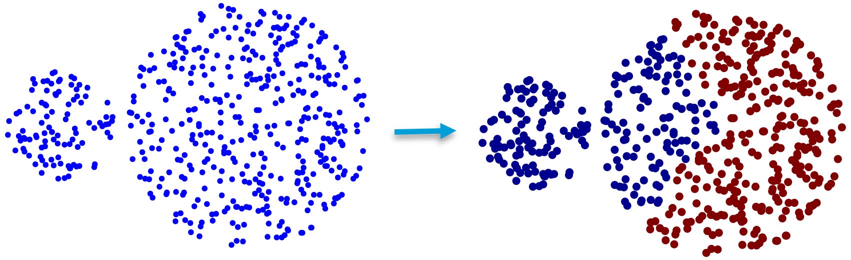

How to Define Inter-Cluster Similarity

- Group Average
  - Middle ground between MIN and MAX.
  - If square distance is used, it is called Ward method.

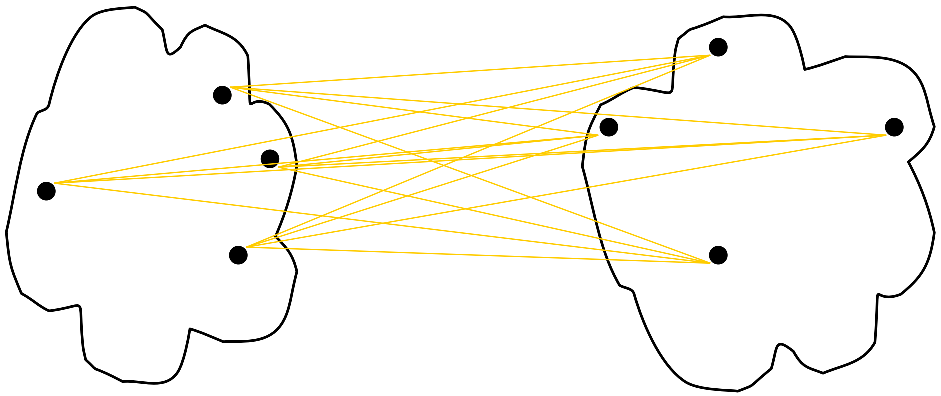

$$\text{proximity}(\text{Cluster}_i, \text{Cluster}_j) = \frac{\sum_{\substack{p_i \in \text{Cluster}_i \\ p_j \in \text{Cluster}_j}} \text{proximity}(p_i, p_j)}{|\text{Cluster}_i| \times |\text{Cluster}_j|}$$

How to Define Inter-Cluster Similarity

- Distance Between Centroids

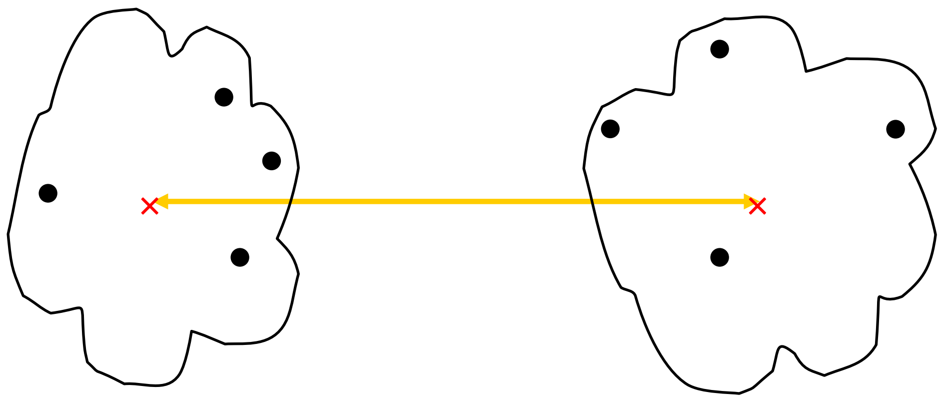

Hierarchical Clustering

Can handle non-elliptical shapes

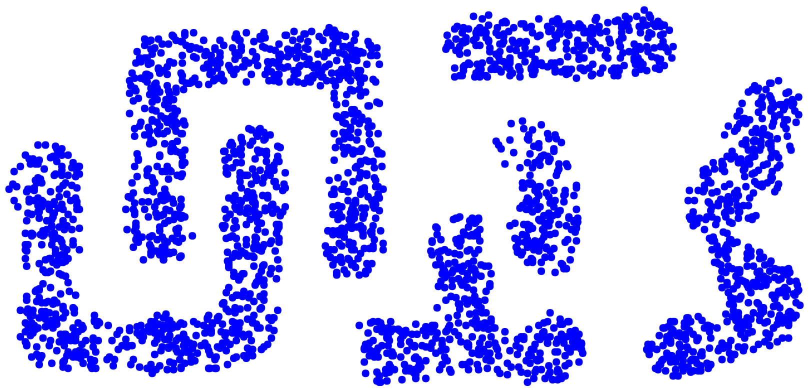

Original Points

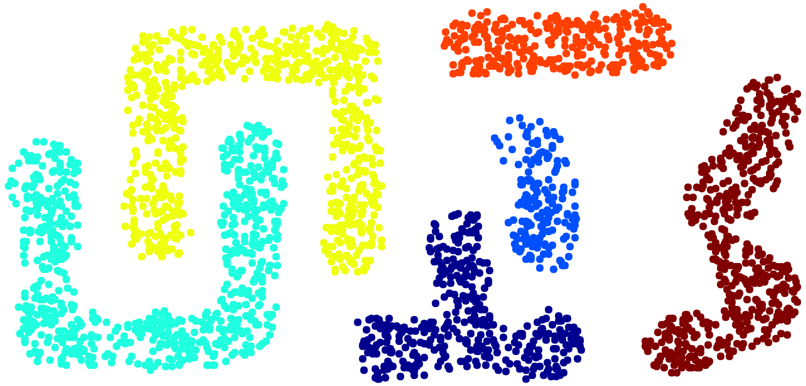

Six Clusters

Hierarchical Clustering

- Visualization + dendograms (easy to find answer when changing number of clusters)
- Hierarchical has a high time complexity (polynomial $O(n^3)$).

Demo time!

Partitional Clustering: K-means

BOSTON UNIVERSITY

K-means Clustering

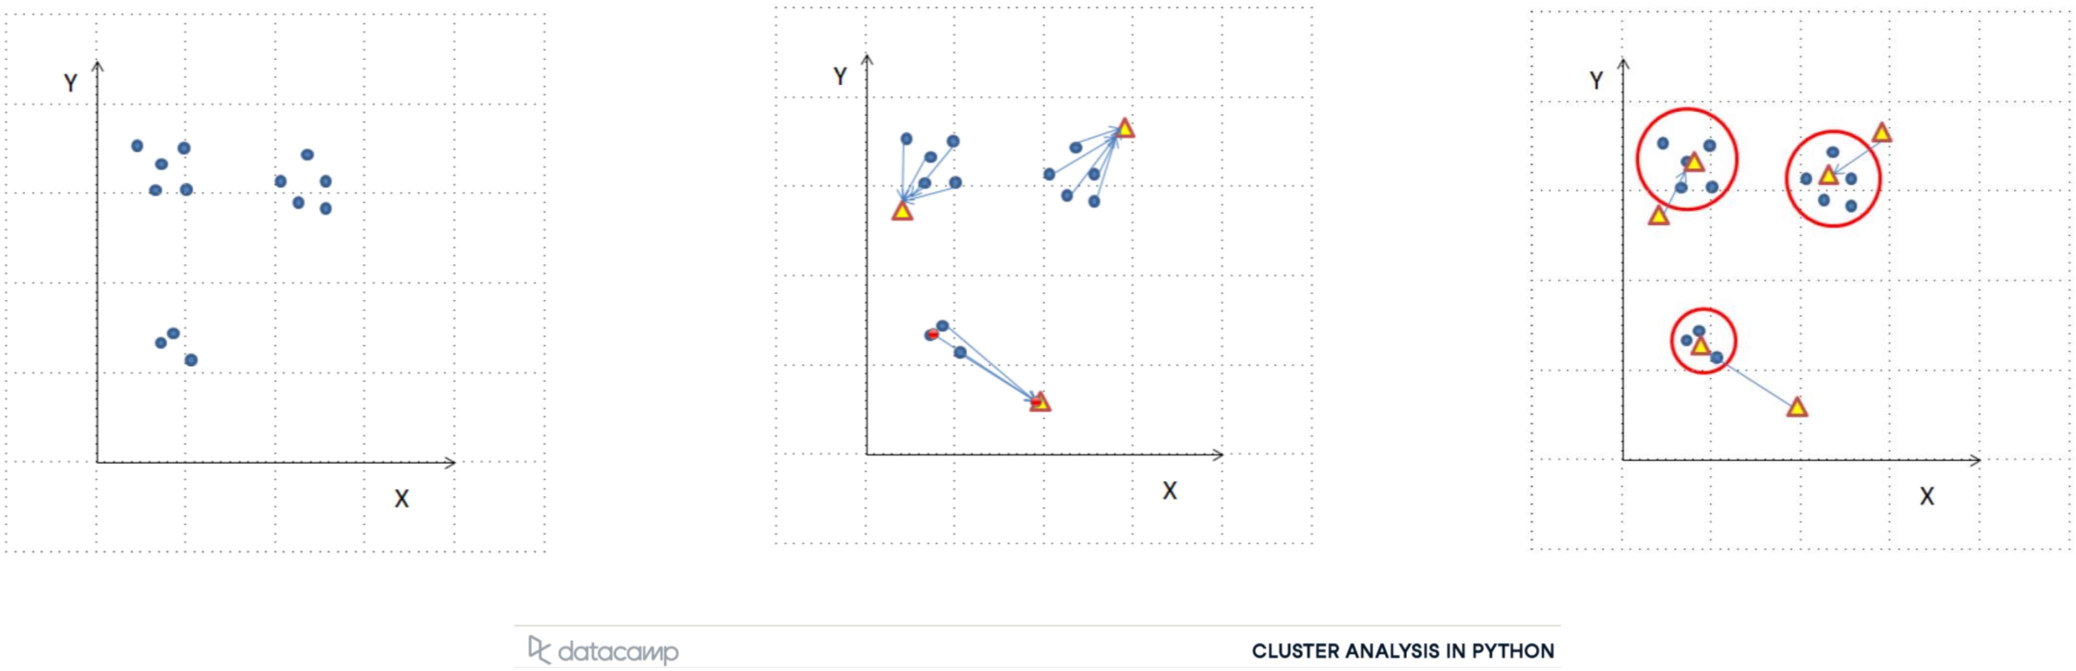

K-means Clustering

- Number of clusters, K, must be specified in advance.
- Each cluster is associated with a *centroid* (center point).
- Each point is assigned to the cluster with the closest centroid.
- The basic algorithm is very simple

1: Select K points as the initial centroids.
2: **repeat**
3: Form K clusters by assigning all points to the closest centroid.
4: Recompute the centroid of each cluster.
5: **until** The centroids don't change

Example of K-means Clustering

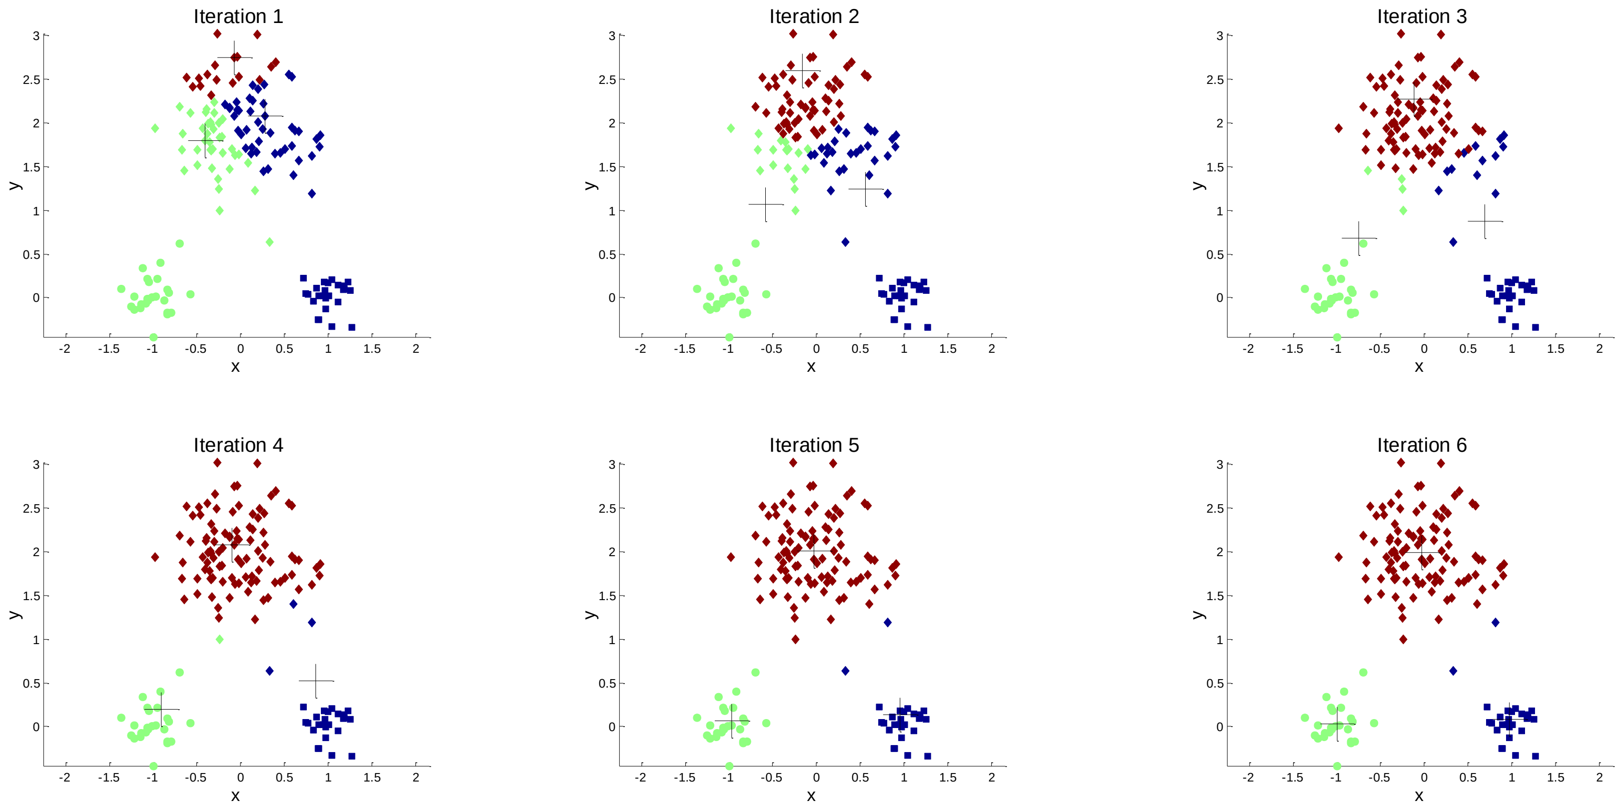

K-means Objective Function

- Euclidean Sum of Squared Error (SSE)
  - For each point, the error is the distance to the nearest cluster center.
  - To get SSE, we square these errors and sum them.

$$SSE = \sum_{i=1}^{K} \sum_{x \in C_i} dist^2(m_i, x)$$

  - This is also called distortion or inertia.
  - x is a data point in cluster C_i and m_i is the centroid (mean) for cluster C_i
  - SSE improves in each iteration of K-means until it reaches a local or global minima.

K-means Playground

- This playground helps understand the mechanism of k-means. Use it to better visualize how the distribution of data affects the clustering.

Effect of Random Initialization

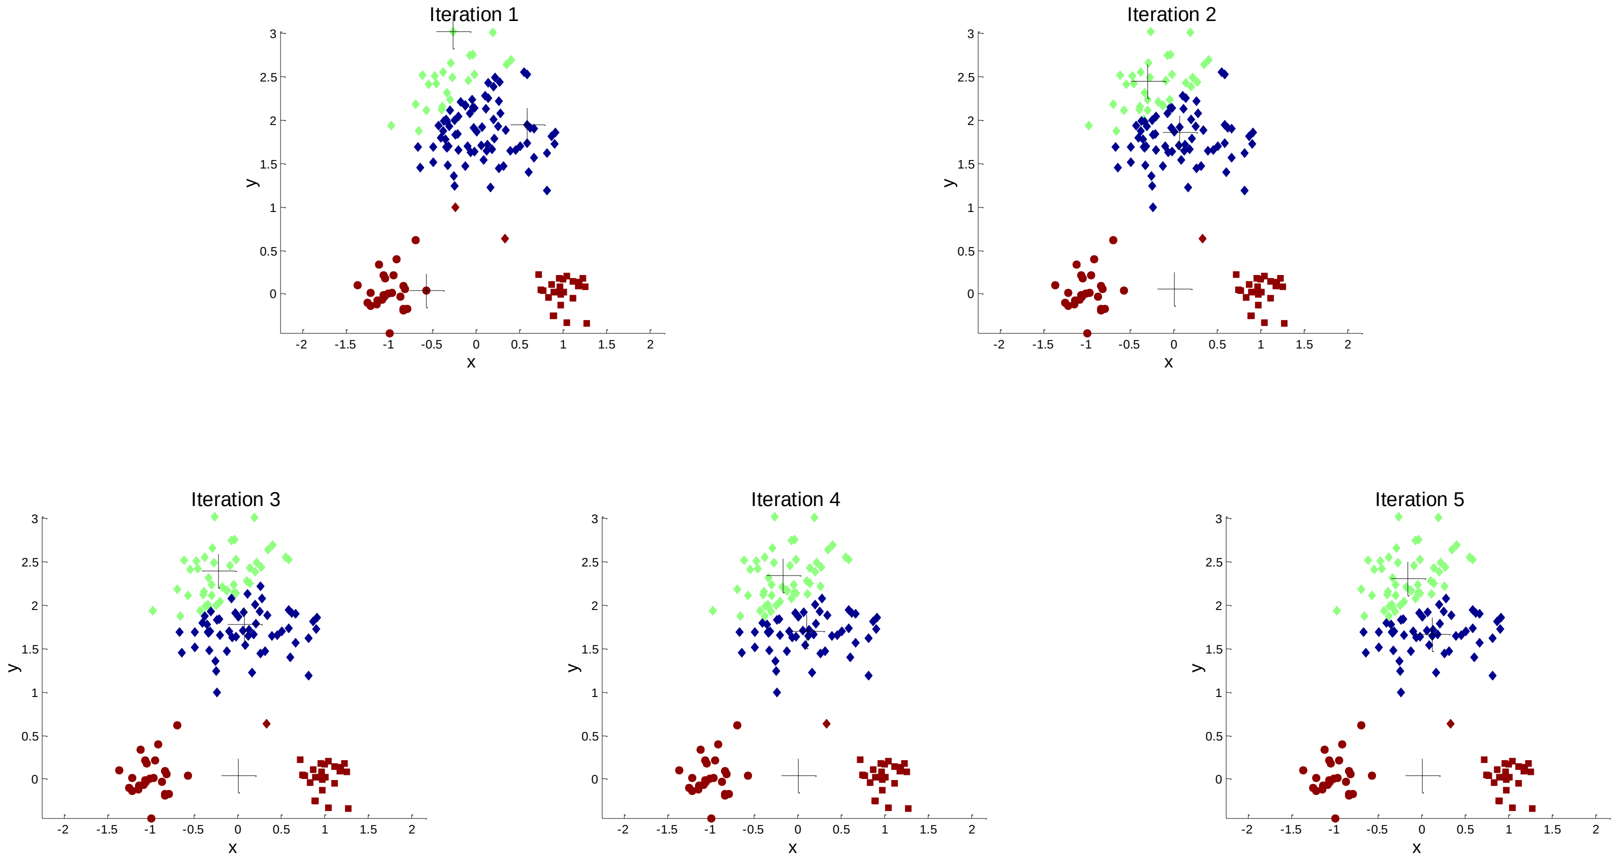

Effect of Random Initialization

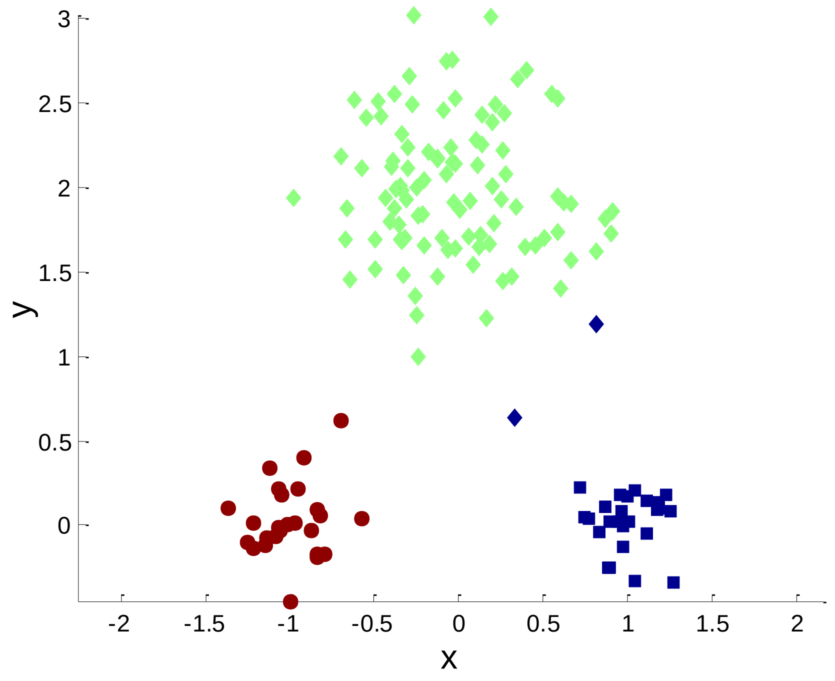

Optimal Clustering

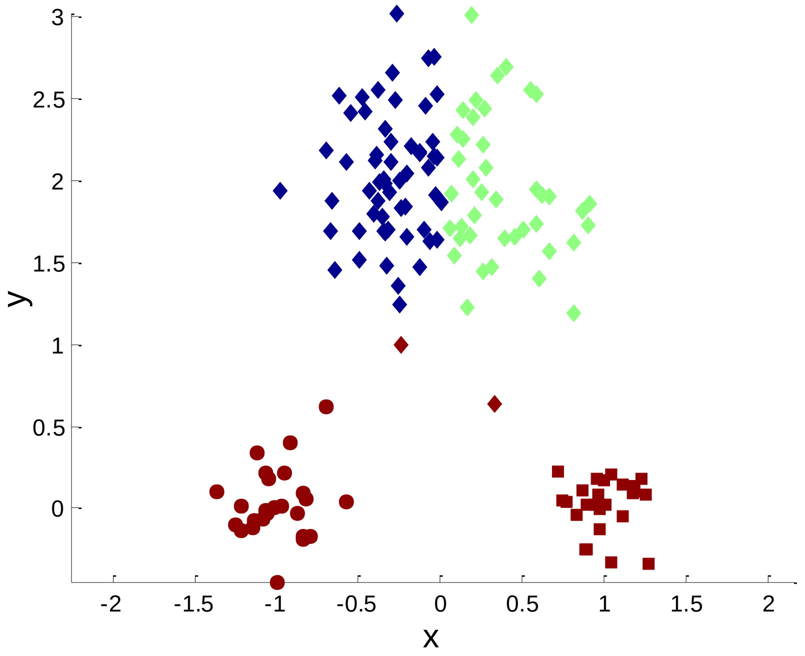

Sub-optimal Clustering

Limitations of K-means

- K-means has problems when clusters are of differing
  - Sizes
  - Densities
  - Non-globular shapes
- K-means is faster than hierarchical clustering.
- K-means is susceptible to suboptimal initialization.
- What do boundaries between clusters look like?
- K-means has problems when the data contains outliers.
  - One possible solution is to remove outliers before clustering

Limitations of K-means: Differing Sizes

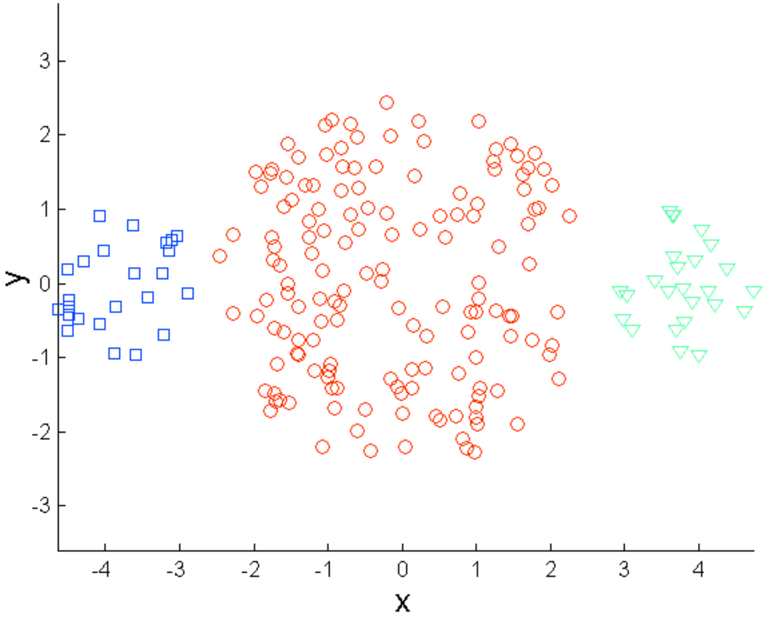

Original Points

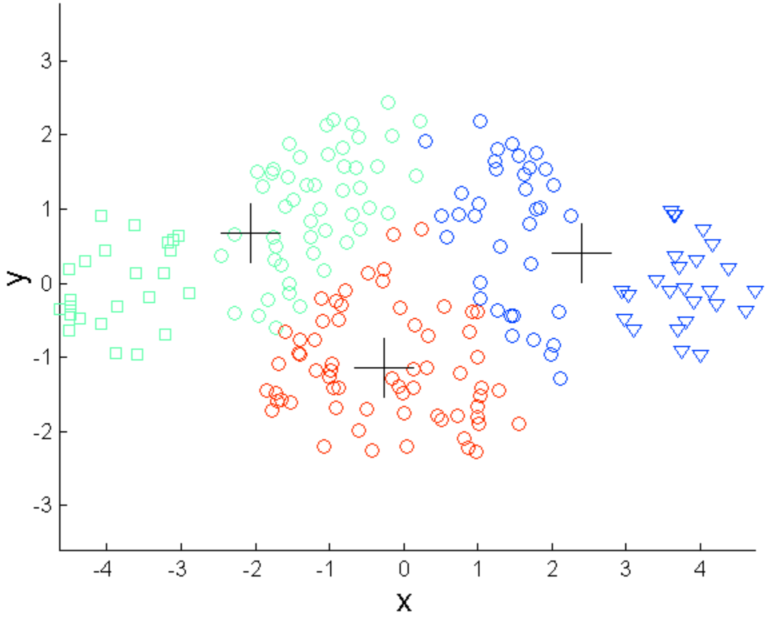

K-means (3 Clusters)

Limitations of K-means: Differing Density

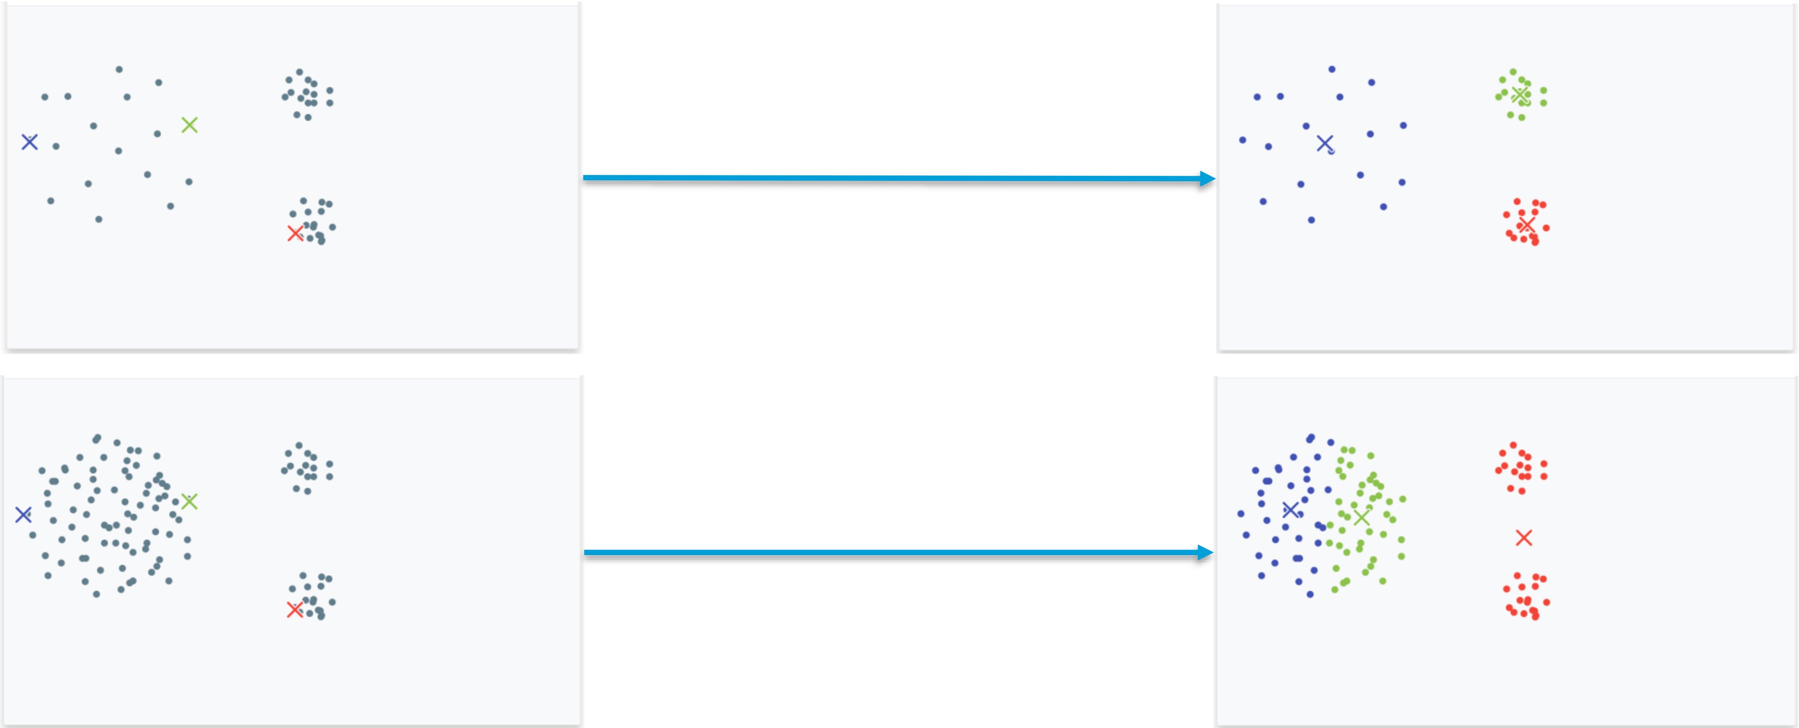

Original Points **K-means (3 Clusters)**

Deciding the number of clusters

Inertia: Sum of squared distances of samples to their respective closest cluster centers.

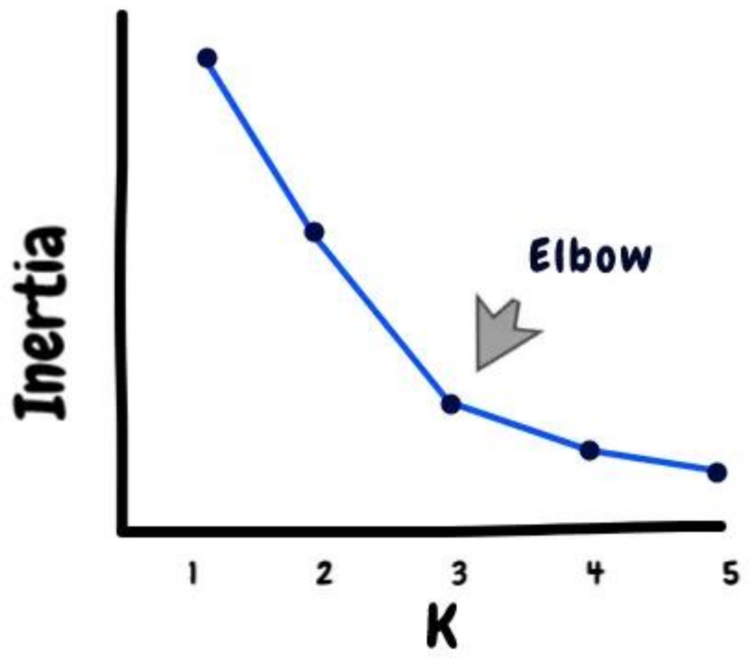

towardsdatascience.com

Measuring Clustering Quality

Method 1: via Labels (Cheating…)

- Testing clustering using supervised learning.
- Calculate the error between ideal labelling and assigned cluster label.

Method 2: via Correlation

- Compute the correlation between the two matrices
- High magnitude of correlation indicates that points that belong to the same cluster are close to each other.
- Not a good measure for hierarchical clustering.
- Similarity can be calculated in many ways.
  - Here, we use $s = \frac{Max - d}{Max}$ or $s = \frac{1}{1+d}$ where *Max* is the maximum possible distance and *d* is the distance measure.

Method 2: via Correlation

- Correlation of ideal similarity and proximity matrices for the K-means clustering

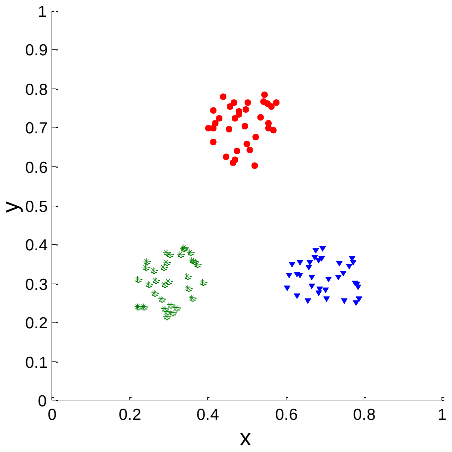

Proximity matrix

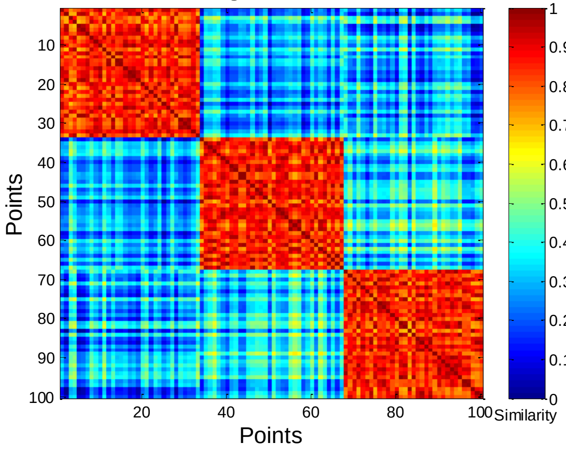

Ideal similarity matrix

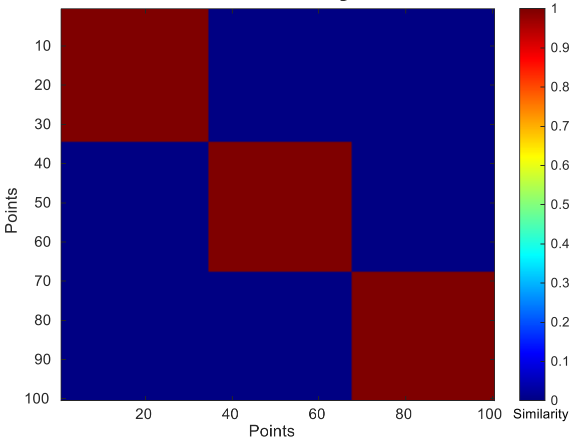

well-clustered data set

Corr = 0.9235

Method 2: via Correlation

- Correlation of ideal similarity and proximity matrices for the K-means clustering

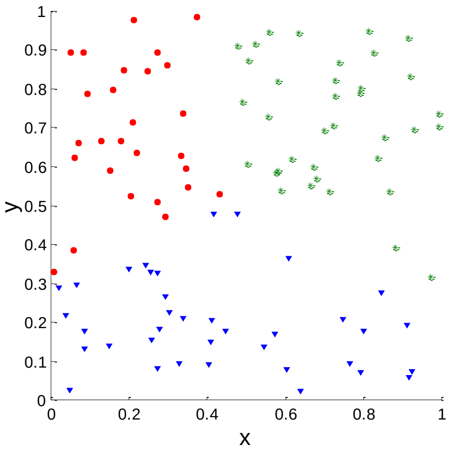

poorly-clustered data set

Proximity matrix

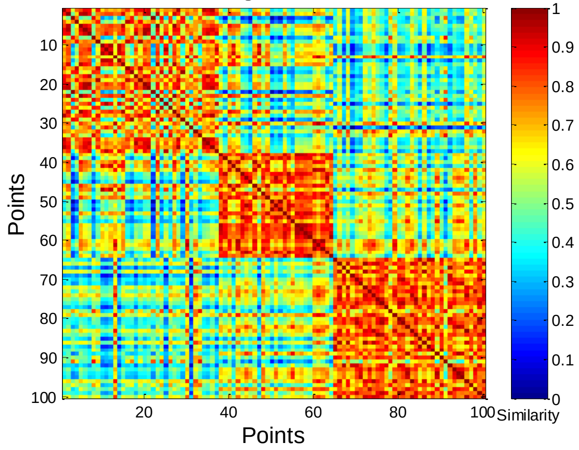

Corr = 0.5810

Ideal similarity matrix